

mlflow

Platform for Machine Learning Lifecycle

Jules S. Damji

@2twitme

Outline – Introduction to MLflow: How to Use MLflow Tracking - Module 1

- Overview of ML development challenges
- How MLflow tackles these
- Concepts and Motivations
- MLFlow Components
 - MLflow Tracking
 - How to use MLflow Tracking APIs
 - Use Databricks Community Edition
 - Explore MLflow UI
 - Tutorials & Exercises
- Q & A

<https://github.com/dmatrix/olt-mlflow>

Machine Learning Development is Complex

Traditional Software vs. Machine Learning

Traditional Software

- **Goal:** Meet a functional specification
- Quality depends only on code
- Typically pick one software stack w/ fewer libraries and tools
- Limited deployment environments

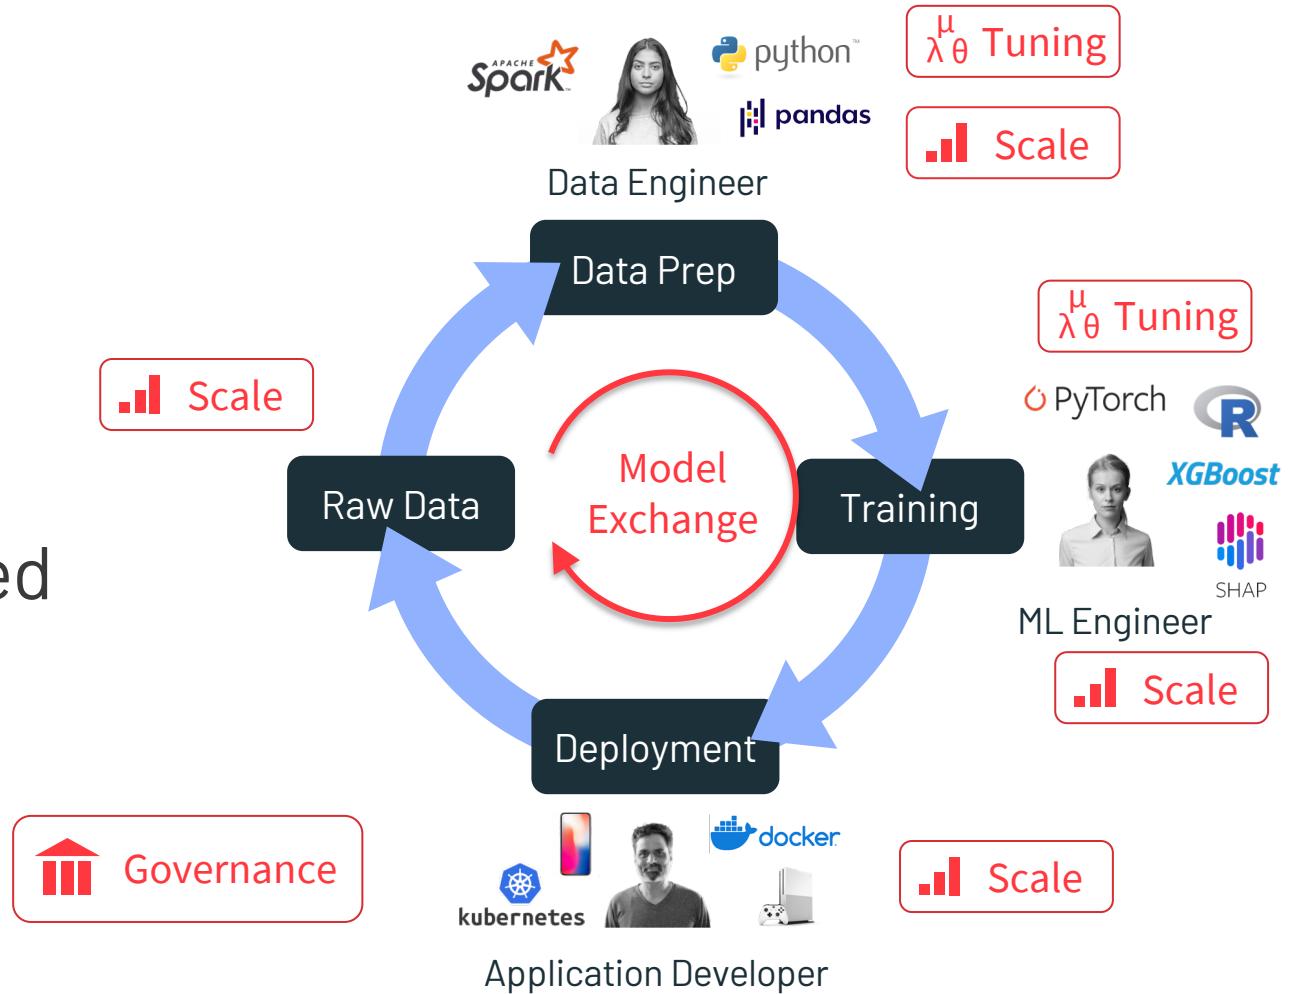
Machine Learning

- **Goal:** Optimize metric (e.g., accuracy). Constantly experiment to improve it
- Quality depends on input data and tuning parameters
- Over time data changes; models drift...
- Compare + combine many libraries, model
- Diverse deployment environments



But Building ML Applications is Complex

- Continuous, iterative process
- Dependent on data
- Many teams and systems involved



Custom Machine Learning Platforms

Some Big Data Companies

- + Standardize the data prep / training / deploy loop:
if you work with the platform, you get these!
- Limited to a few algorithms or frameworks
- Tied to one company's infrastructure
- Out of luck if you left the company....

Can we provide similar benefits in an **open** manner?

mlflow: An Open Source ML Platform

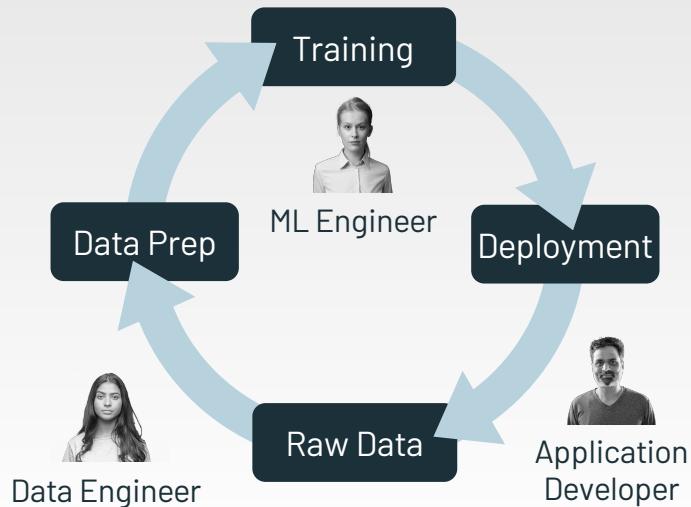
mlflow
TRACKING
Experiment management

mlflow
PROJECTS
Reproducible runs

mlflow
MODELS
Model packaging and deployment

mlflow
MODEL REGISTRY
Model management

Any Language



Any ML Library



MLflow Design Philosophy

API-First

- Submit runs, log models, metrics, etc. from popular library & language
- Abstract “model” lambda function that MLflow can then deploy in many places (Docker, Azure ML, Spark UDF)
- Open interface allows easy integration from the community

**Key enabler: built around
Programmatic APIs, REST APIs & CLI**

Modular Design

- Allow different components individually (e.g., use MLflow’s project format but not its deployment tools)
- Not monolithic
- But Distinctive and Selective

**Key enabler: distinct components
(Tracking/Projects/Models/Registry)**

Introducing **mlflow**

Open machine learning platform

Works with popular ML library & language

Runs the same way anywhere (e.g., any cloud or locally)

Designed to be useful for 1 or 1000+ person orgs

Simple. Modular. Easy-to-use.

Offers positive developer experience to get started!

Key Concepts in MLflow Tracking

Parameters: key-value inputs to your code

Metrics: numeric values (can update over time)

Tags and Notes: information about a run

Artifacts: files, data, and models

Source: what code ran?

Version: what of the code?

Run: an instance of code that runs by MLflow

Experiment: {Run, ... Run}

Model Development without MLflow Tracking

```
data    = load_text(file_name=file)
ngrams = extract_ngrams(data, N=n)
model   = train_model(ngrams,
                      learning_rate=lr)
score   = compute_accuracy(model)

print("For n=%d, lr=%f: accuracy=%f"
      % (n, lr, score))

pickle.dump(model, open("model.pkl"))
```

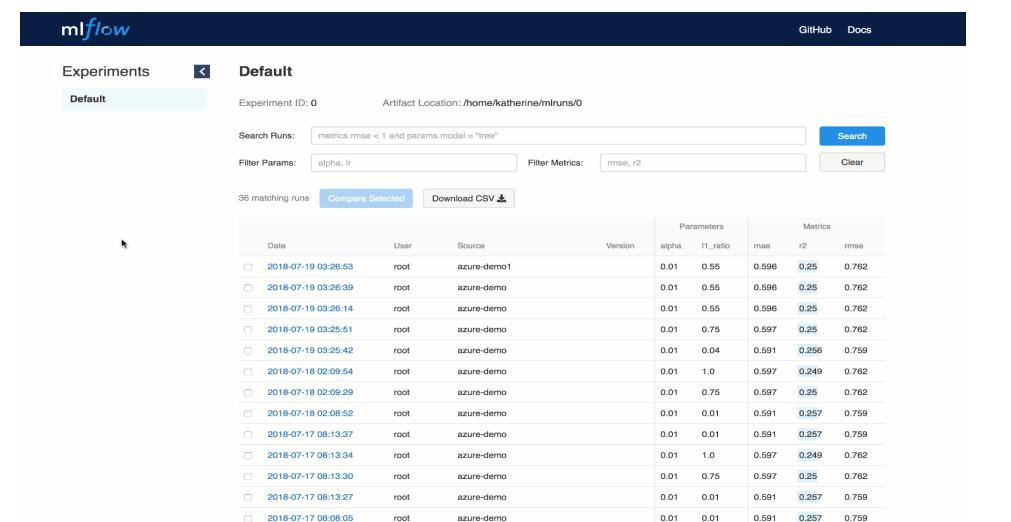
```
For n=2, lr=0.1: accuracy=0.71
For n=2, lr=0.2: accuracy=0.79
For n=2, lr=0.5: accuracy=0.83
For n=2, lr=0.9: accuracy=0.79
For n=3, lr=0.1: accuracy=0.83
For n=3, lr=0.2: accuracy=0.82
For n=4, lr=0.5: accuracy=0.75
...
```

What version of
my code was this
result from?

Model Development with MLflow is Simple!

```
import mlflow
data    = load_text(file_name=file)
ngrams = extract_ngrams(data, N=n)
model   = train_model(ngrams,
                      learning_rate=lr)
score   = compute_accuracy(model)
with mlflow.start_run():
    mlflow.log_param("data_file", file)
    mlflow.log_param("n", n)
    mlflow.log_param("learn_rate", lr)
    mlflow.log_metric("score", score)
    mlflow.sklearn.log_model(model)
```

\$ mlflow ui



The screenshot shows the MLflow UI interface. At the top, there's a search bar with the query "metrics.rmse < 1 and params.model = 'tree'". Below the search bar, there are sections for "Experiment ID: 0" and "Artifact Location: /home/katherine/mlruns/0". A table below lists 36 matching runs, each with columns for Date, User, Source, Version, Parameters (alpha, H1_ratio), and Metrics (rmse, r2). The table shows multiple runs with different parameter values and metric scores.

Date	User	Source	Version	Parameters	Metrics
2018-07-19 03:26:53	root	azure-demo1	0.01	0.55 0.596 0.25	0.596 0.25 0.762
2018-07-19 03:26:39	root	azure-demo	0.01	0.55 0.596 0.25	0.596 0.25 0.762
2018-07-19 03:26:14	root	azure-demo	0.01	0.55 0.596 0.25	0.596 0.25 0.762
2018-07-19 03:25:51	root	azure-demo	0.01	0.75 0.597 0.25	0.597 0.25 0.762
2018-07-19 03:25:42	root	azure-demo	0.01	0.04 0.591 0.256	0.591 0.256 0.759
2018-07-18 02:09:54	root	azure-demo	0.01	1.0 0.597 0.249	0.597 0.249 0.762
2018-07-18 02:09:29	root	azure-demo	0.01	0.75 0.597 0.25	0.597 0.25 0.762
2018-07-18 02:08:52	root	azure-demo	0.01	0.01 0.591 0.257	0.591 0.257 0.759
2018-07-17 08:13:37	root	azure-demo	0.01	0.01 0.591 0.257	0.591 0.257 0.759
2018-07-17 08:13:34	root	azure-demo	0.01	1.0 0.597 0.249	0.597 0.249 0.762
2018-07-17 08:13:30	root	azure-demo	0.01	0.75 0.597 0.25	0.597 0.25 0.762
2018-07-17 08:13:27	root	azure-demo	0.01	0.01 0.591 0.257	0.591 0.257 0.759
2018-07-17 08:08:05	root	azure-demo	0.01	0.01 0.591 0.257	0.591 0.257 0.759

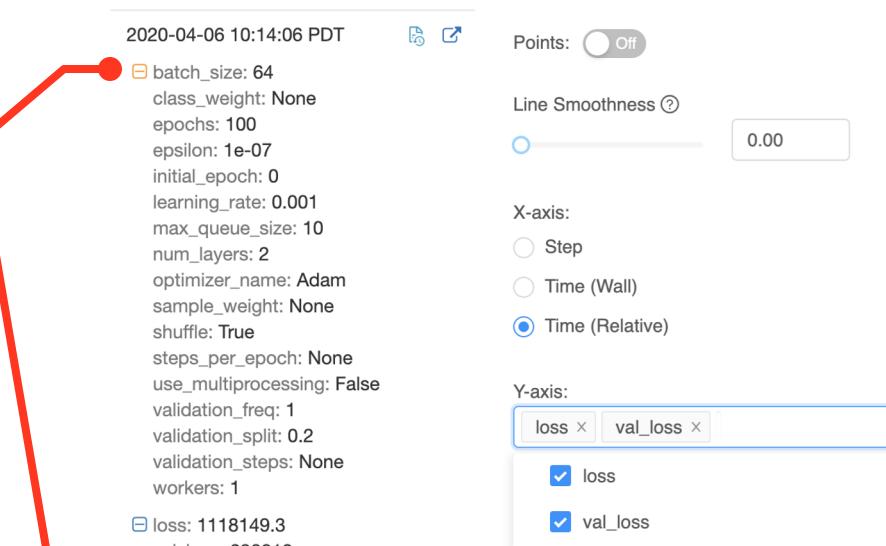
Track parameters, metrics, artifacts, output files & code version

mlflow Tracking for ML Experiments

Easily track parameters, metrics, and artifacts in popular ML libraries

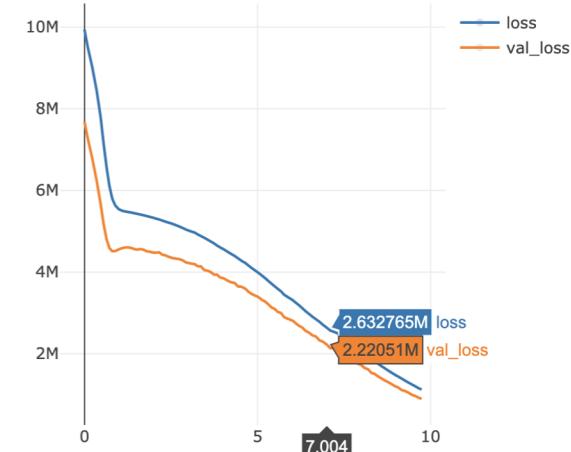
```
with mlflow.start_run(run_name='keras'):
    # log model and datasource
    mlflow.keras.autolog()
```

Library integrations:

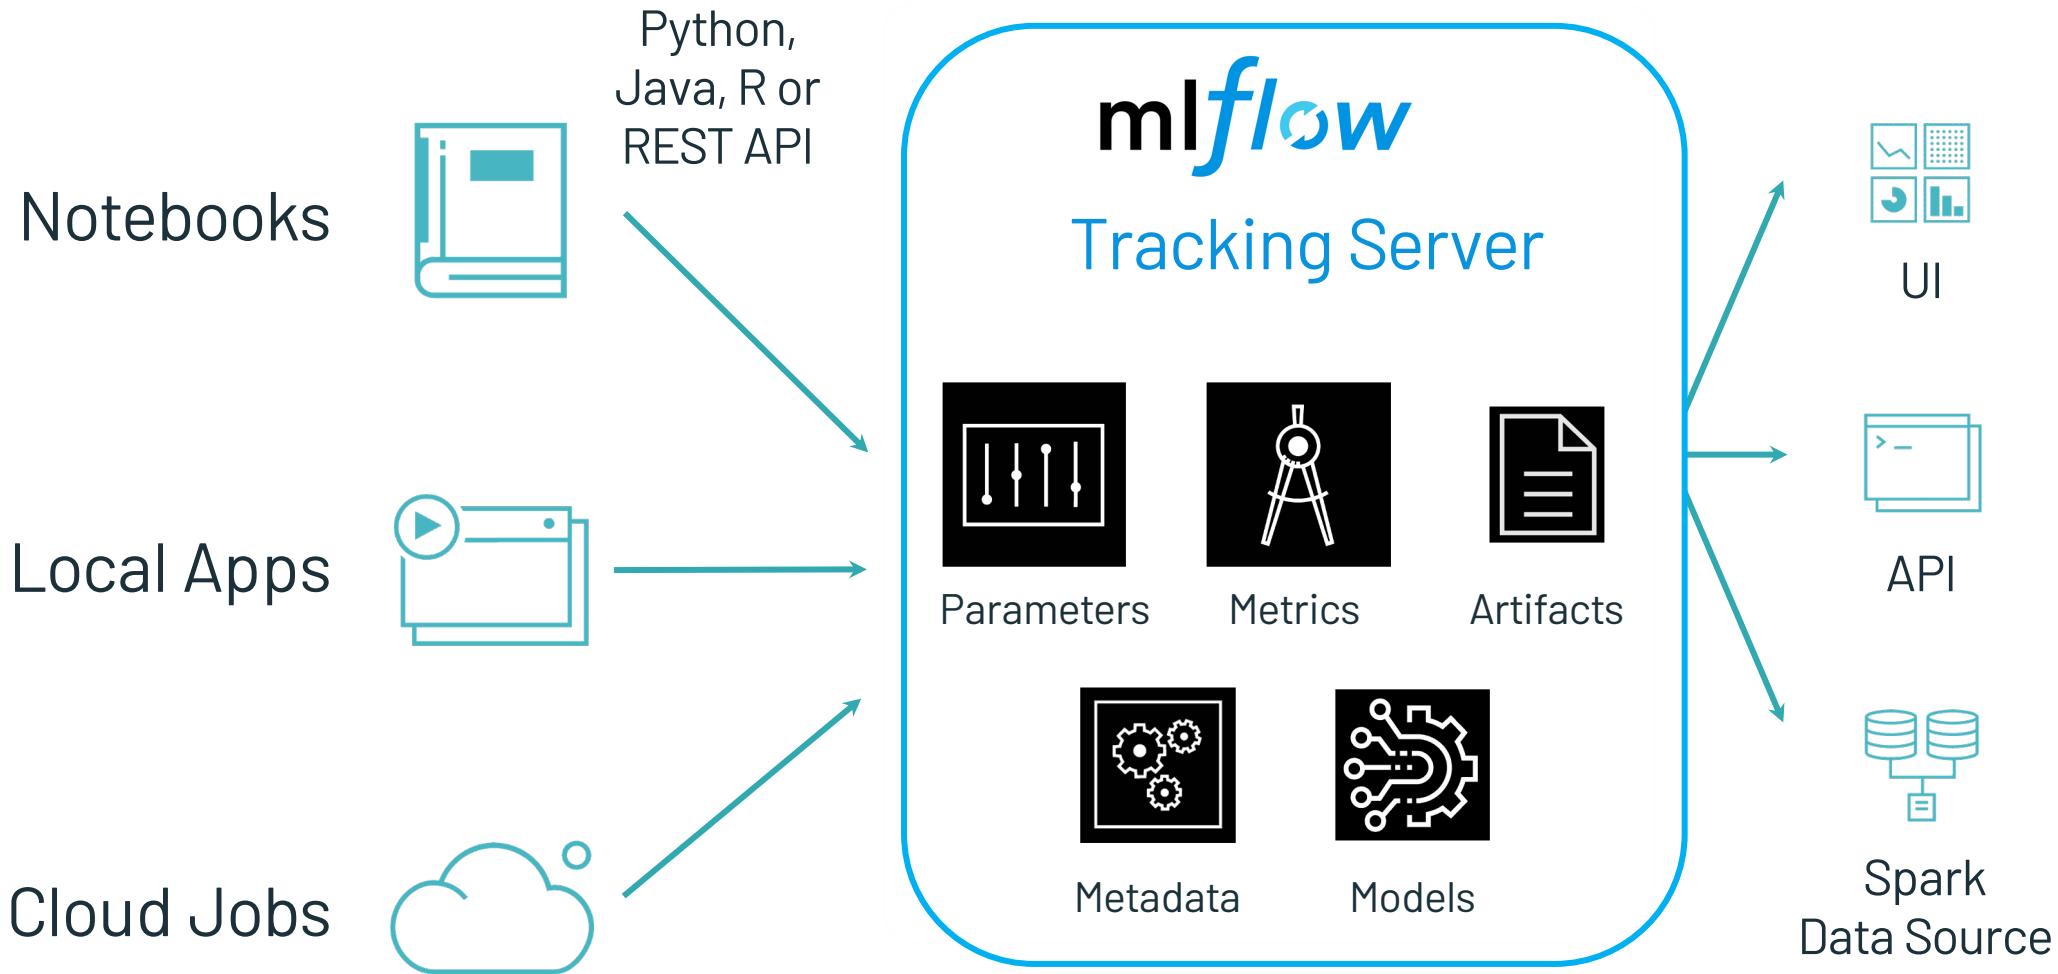


Full Path: dbfs:/databricks/mlflow/25120103/c39...
Size: 303B

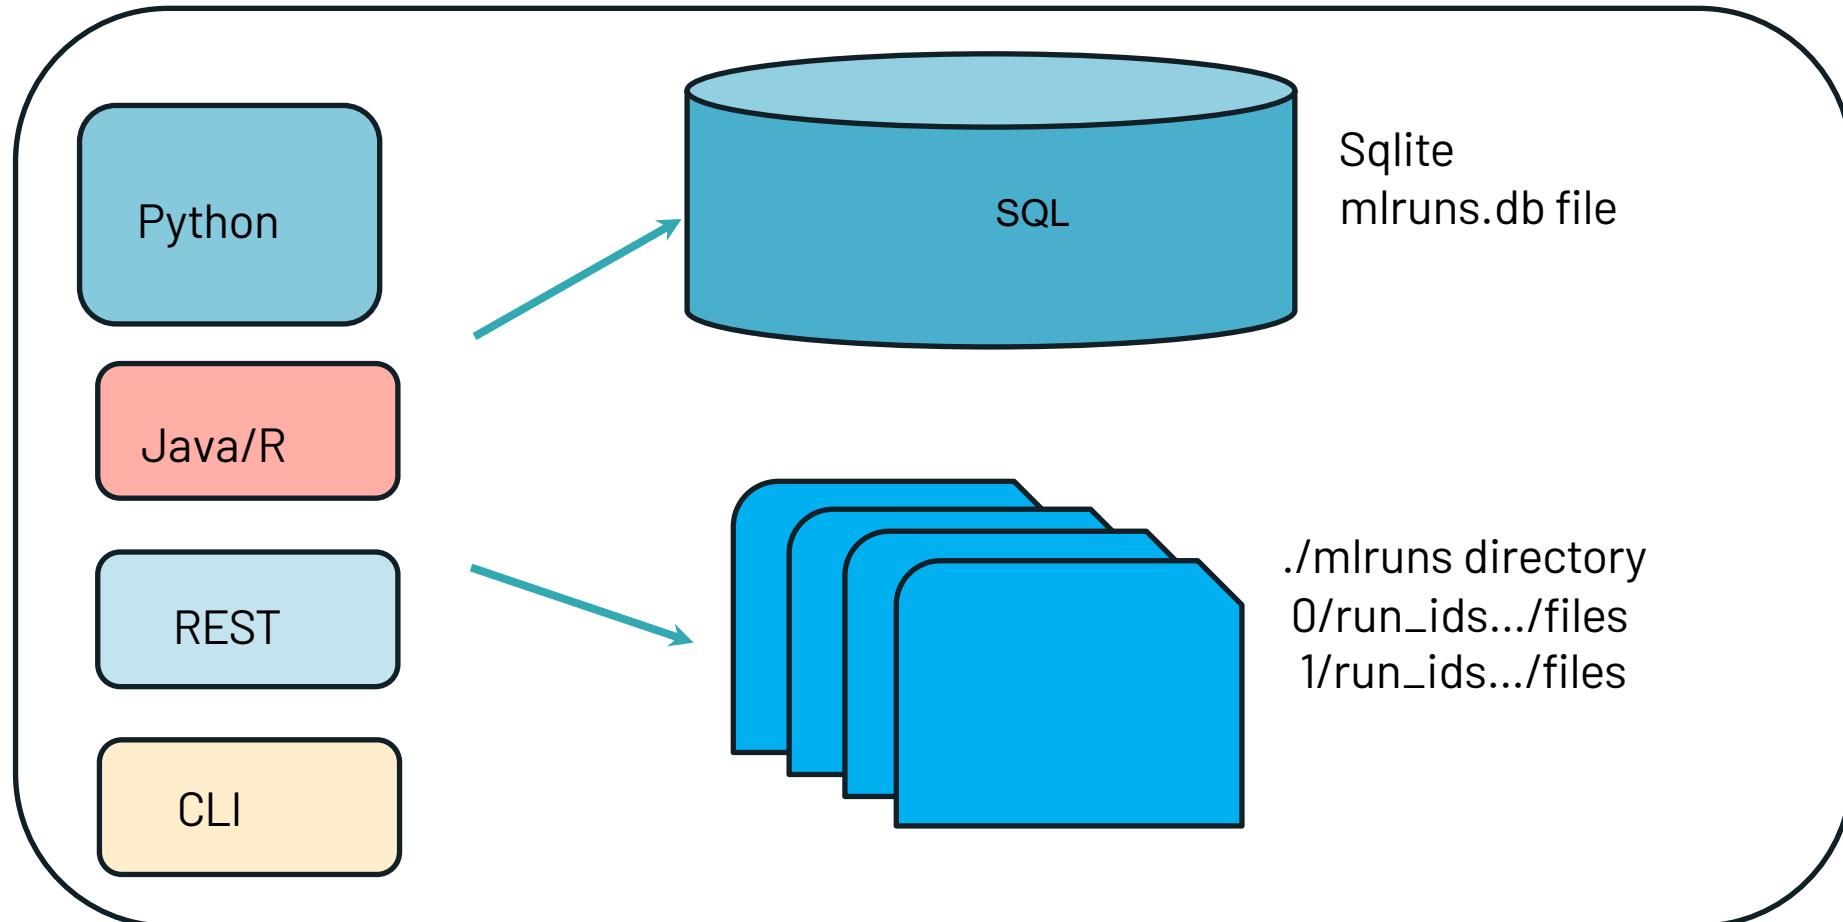
artifact_path: model
flavors:
keras:
data: data
keras_module: keras
keras_version: 2.2.5
python_function:
data: data
env: conda.yaml
loader_module: mlflow.keras
python_version: 3.7.3
run_id: c3995babe1754238acd7b241cad12d7c
utc_time_created: '2020-04-06 17:14:18.343005'

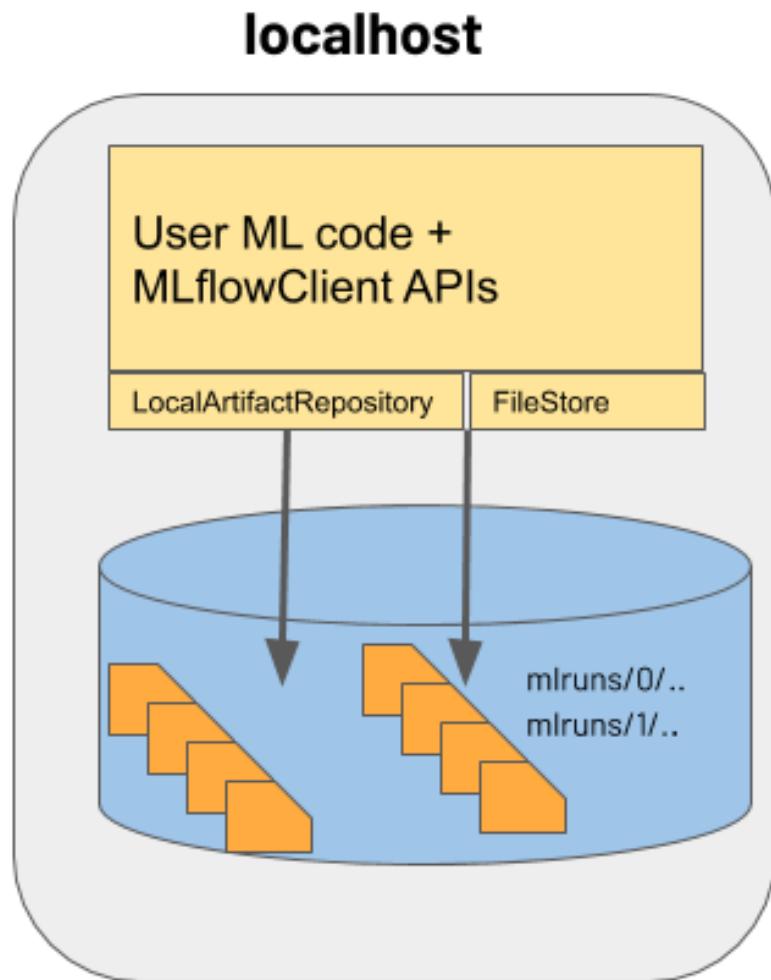


MLflow Tracking Server



MLflow Tracking Local Storage



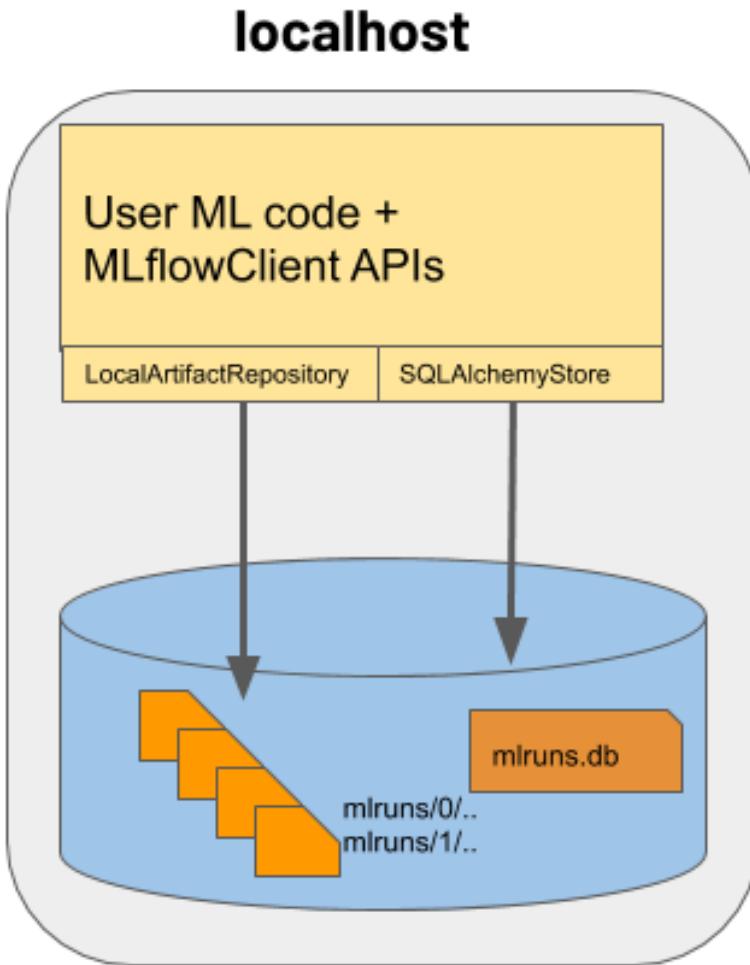


Scenario 1: MLflow on the localhost

No tracking server is involved here, as we run everything locally on our localhost or laptop, hence no REST APIs. The **MLflowClient** directly interfaces with the **FileStore**. This store is used both as an artifact store and backend store under the `./mlruns` directory.

Interaction and Flow:

1. **MLflowClient** --> creates an instance of **LocalArtifactRepository** (to store artifacts)
2. **MLflowClient** --> creates an instance of **FileStore** (to save MLflow entities)
3. Both artifacts and MLflow entities are stored on the local disk under `./mlruns/....`



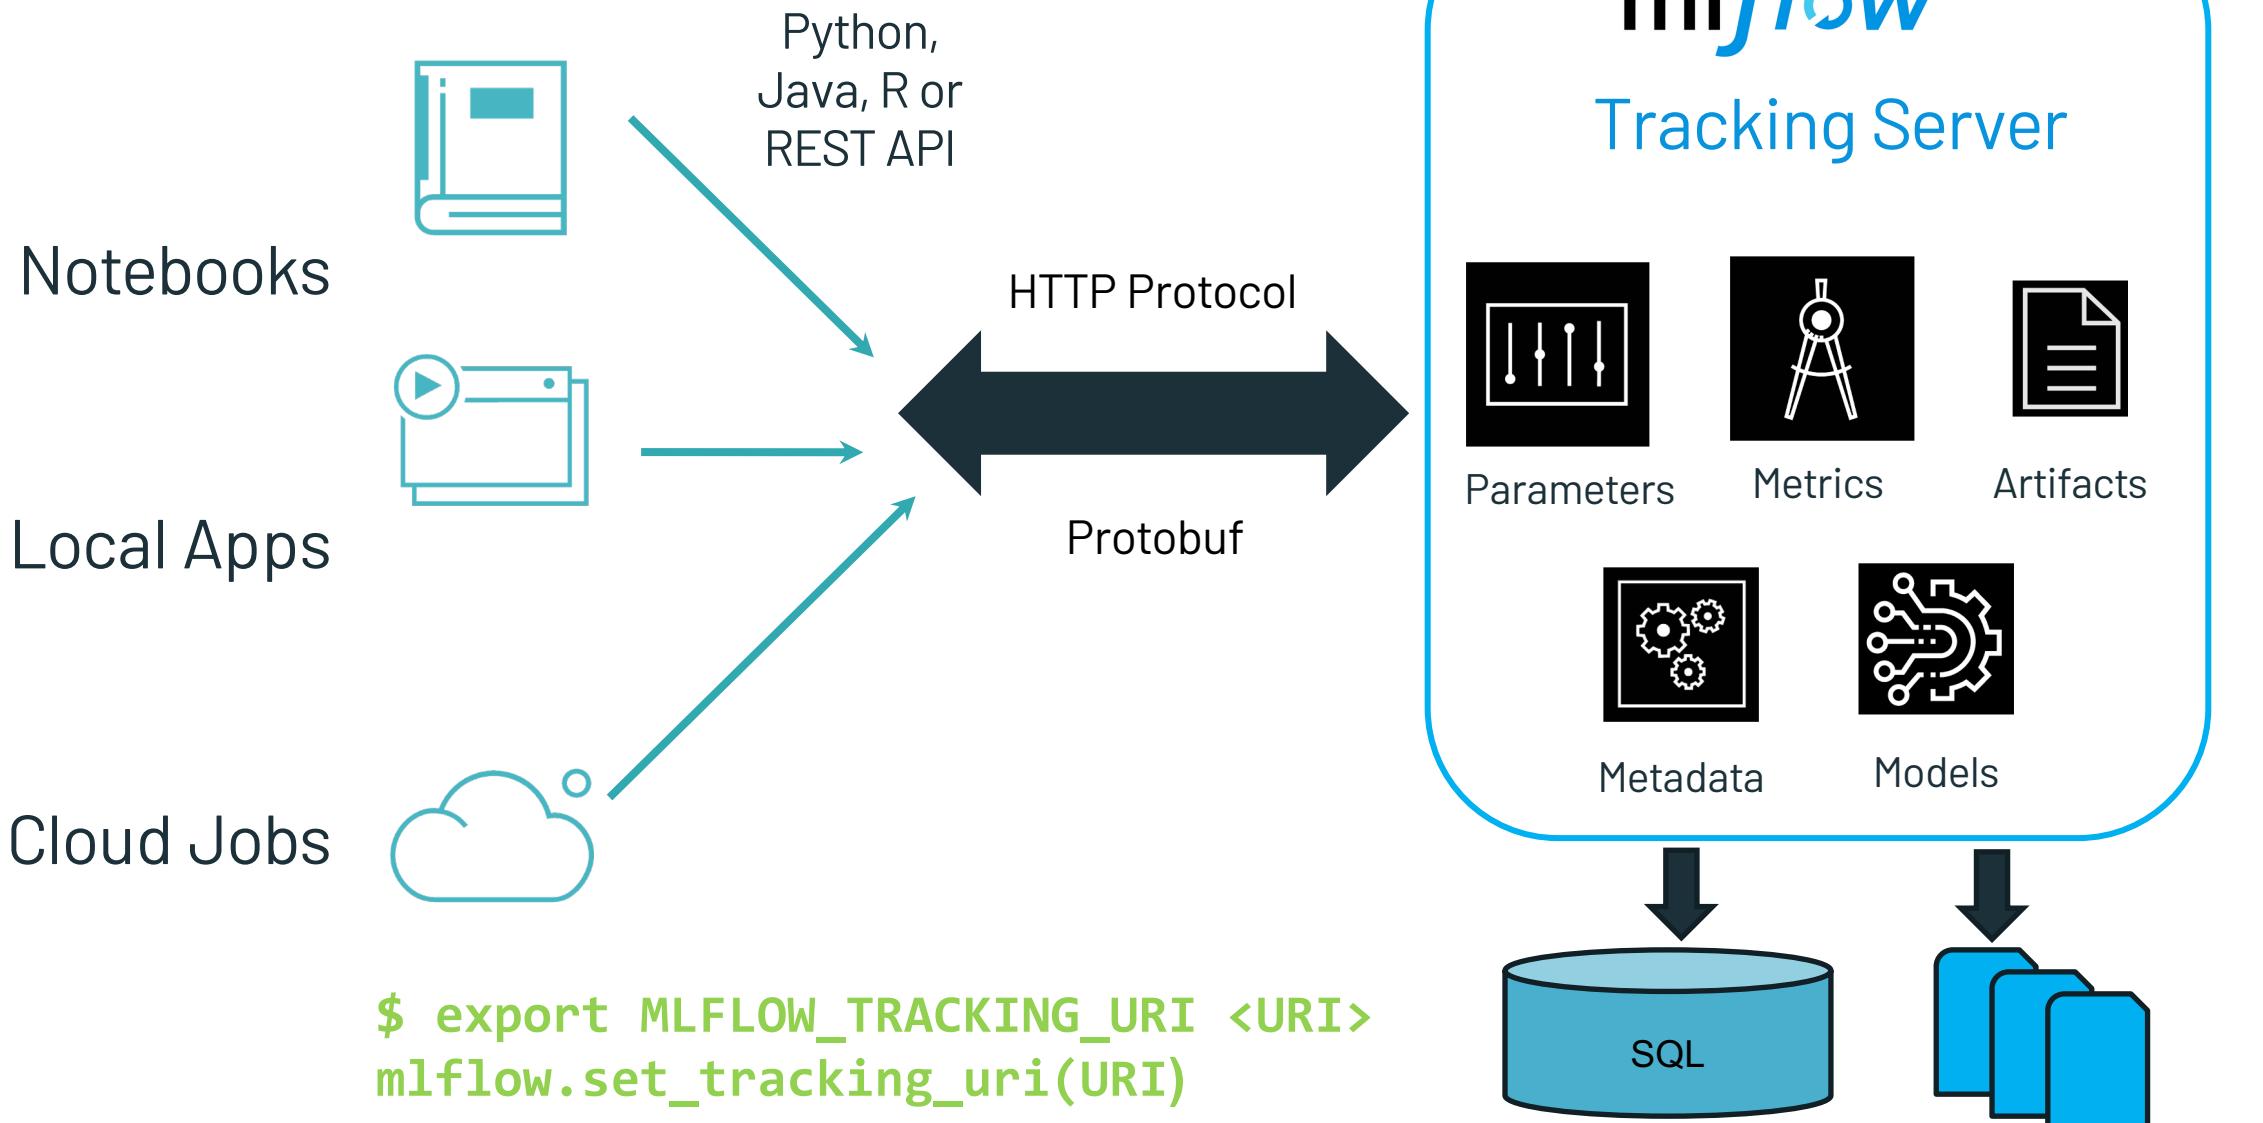
Scenario 2: MLflow on the localhost
with backend store as an
SQLAlchemy compatible database
type: sqlite

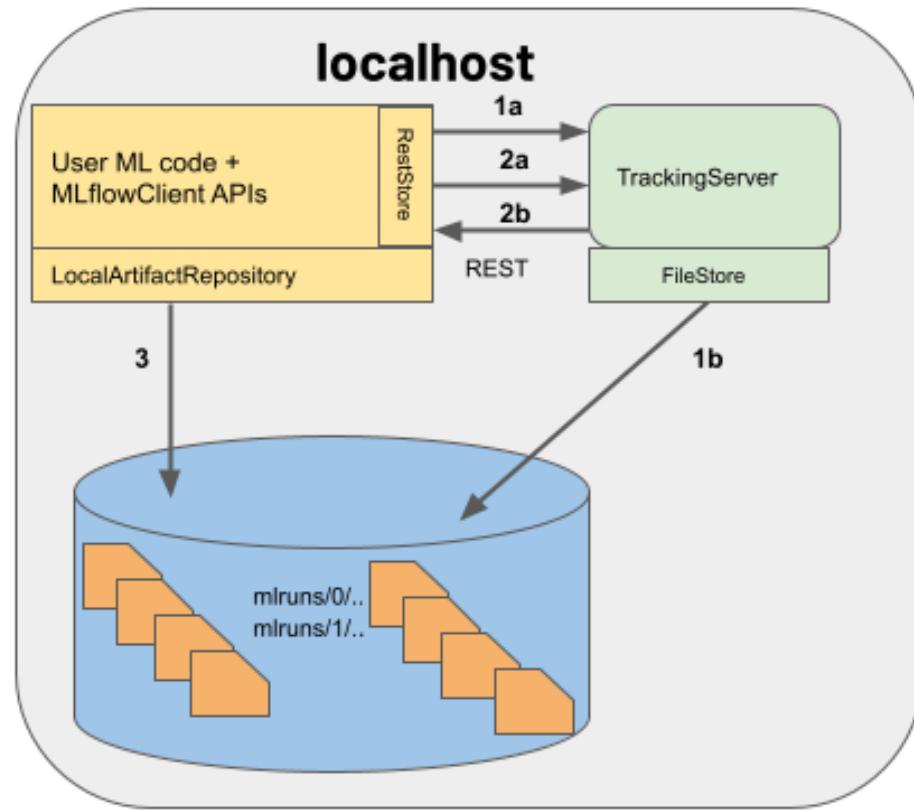
No tracking server, but use an SQLAlchemy compatible
backend store for MLflow entities. Artifacts are stored
locally in the local `./mlruns` directory, and MLflow
entities are stored in an sqlite database file `mlruns.db`

Interaction and Flow:

1. **MLflowClient** --> creates an instance of
LocalArtifactRepository (to save artifacts)
2. **MLflowClient** --> creates an instance of
SQLAlchemyStore (to save MLflow entities)
and writes to sqlite file **mlruns.db**)

MLflow Tracking Server





Scenario 3: Tracking server launched at localhost (default port 5000) : `mlflow server --backend-store-uri file://path/mlruns`

1. Backend store and Artifact Store both use local `file://path/mlruns`. That is, they use the **LocalArtifactRepository** to save artifacts and **FileStore** to save MLflow entities (runs, params, metrics, tags, etc).
2. Client will use an instance of **RestStore** and make REST APIs calls to the tracking server at default port 5000

Interaction and Flow:

MLflow Entities:

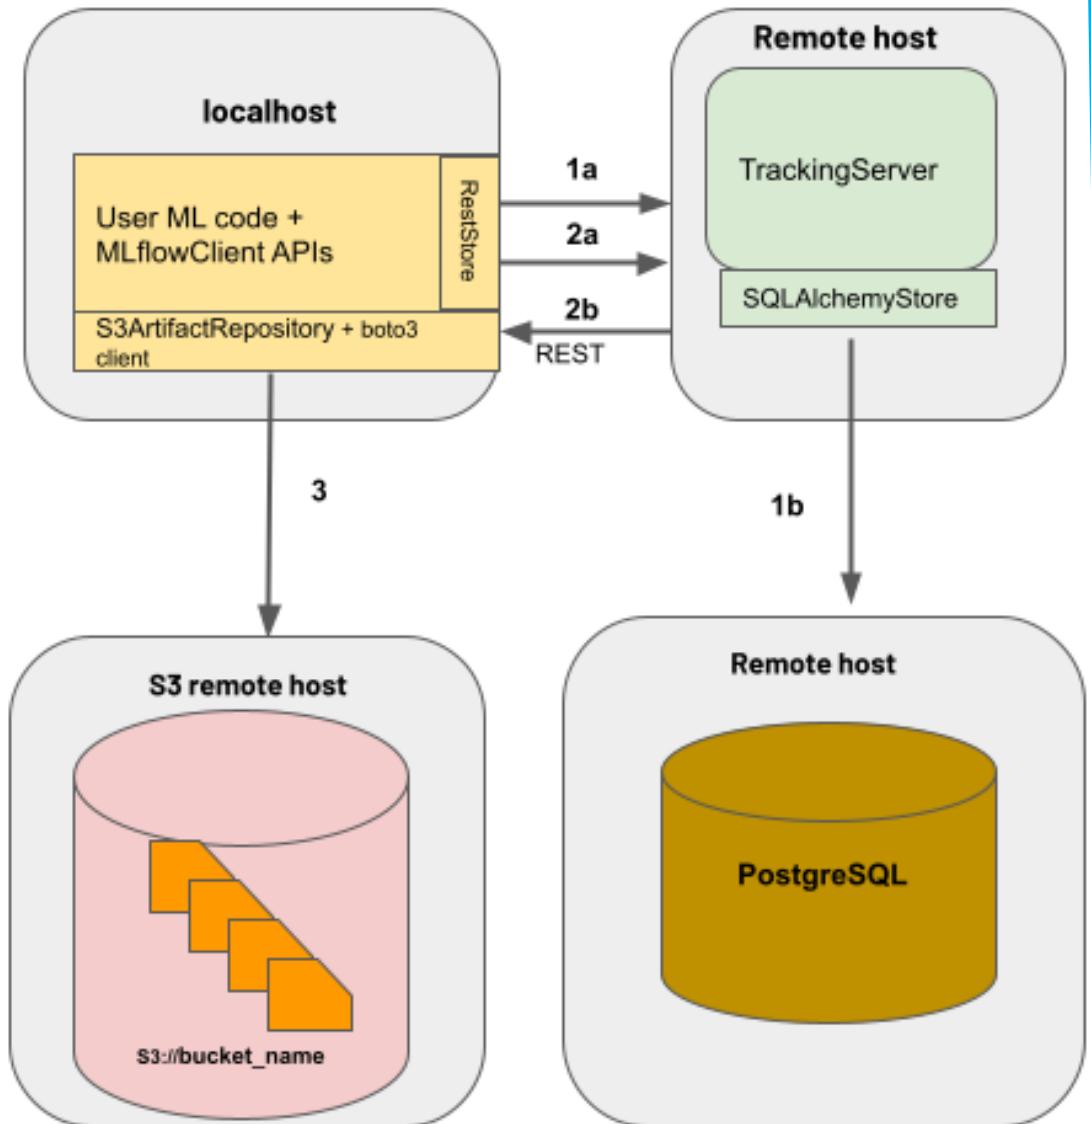
part 1a & b: **MLflowClient** --> creates an instance of **RestStore** --> REST Log MLflow entities Request API Call --> **Tracking Server** --> instance of **FileStore** (to save MLflow entities—params, runs, metrics, etc)

Artifacts:

part 2a: **MLflowClient** --> **RestStore** --> REST Request API Call --> **Tracking Server** (fetch artifact store URI)

part 2b: **Tracking Server** --> REST Response with artifact store URI --> **MLflowClient**

part 3: **MLflowClient** --> creates an instance of **LocalArtifactRepository** (to stores artifacts)



Scenario 4: `mlflow server --backend-store-uri postgresql://URI --default-artifact-root S3:/bucket_name --host hostname`

1. Backend store is a PostgreSQL://URI at a remote host
2. URI scheme-based concrete class of **S3ArtifactRepository** for URI S3:/bucket_name

MLflow Entities:

part 1a & b: **MLflowClient** --> **RestStore** --> **REST Request API Call** --> **Tracking Server** --> creates an instance of **SQLAlchemyStore** (to store MLflow entities, params, runs, metrics, etc) connects to remote host Postgres DB

Artifacts:

part 2a: **MLflowClient** --> **RestStore** --> **REST Request API Call** --> **Tracking Server** (fetch artifact store URI)

part 2b: **Tracking Server** --> **REST Response Response with artifact store URI** --> **MLflowClient**

part 3: **MLflowClient** --> instance of **S3ArtifactRepository** --> S3 remote host bucket (to store artifacts)

The **MLflowClient** will use the **S3ArtifactRepository** (`s3:/bucket_name/`) to save artifacts on the remote S3 bucket, and the **Tracking Server** will use an existing instance of **SQLAlchemyStore** (`postgresql://URI`) to save MLflow entities (runs, params, metrics, tags, etc), after each REST request to log MLflow entities.

MLflow Tracking Server Storage

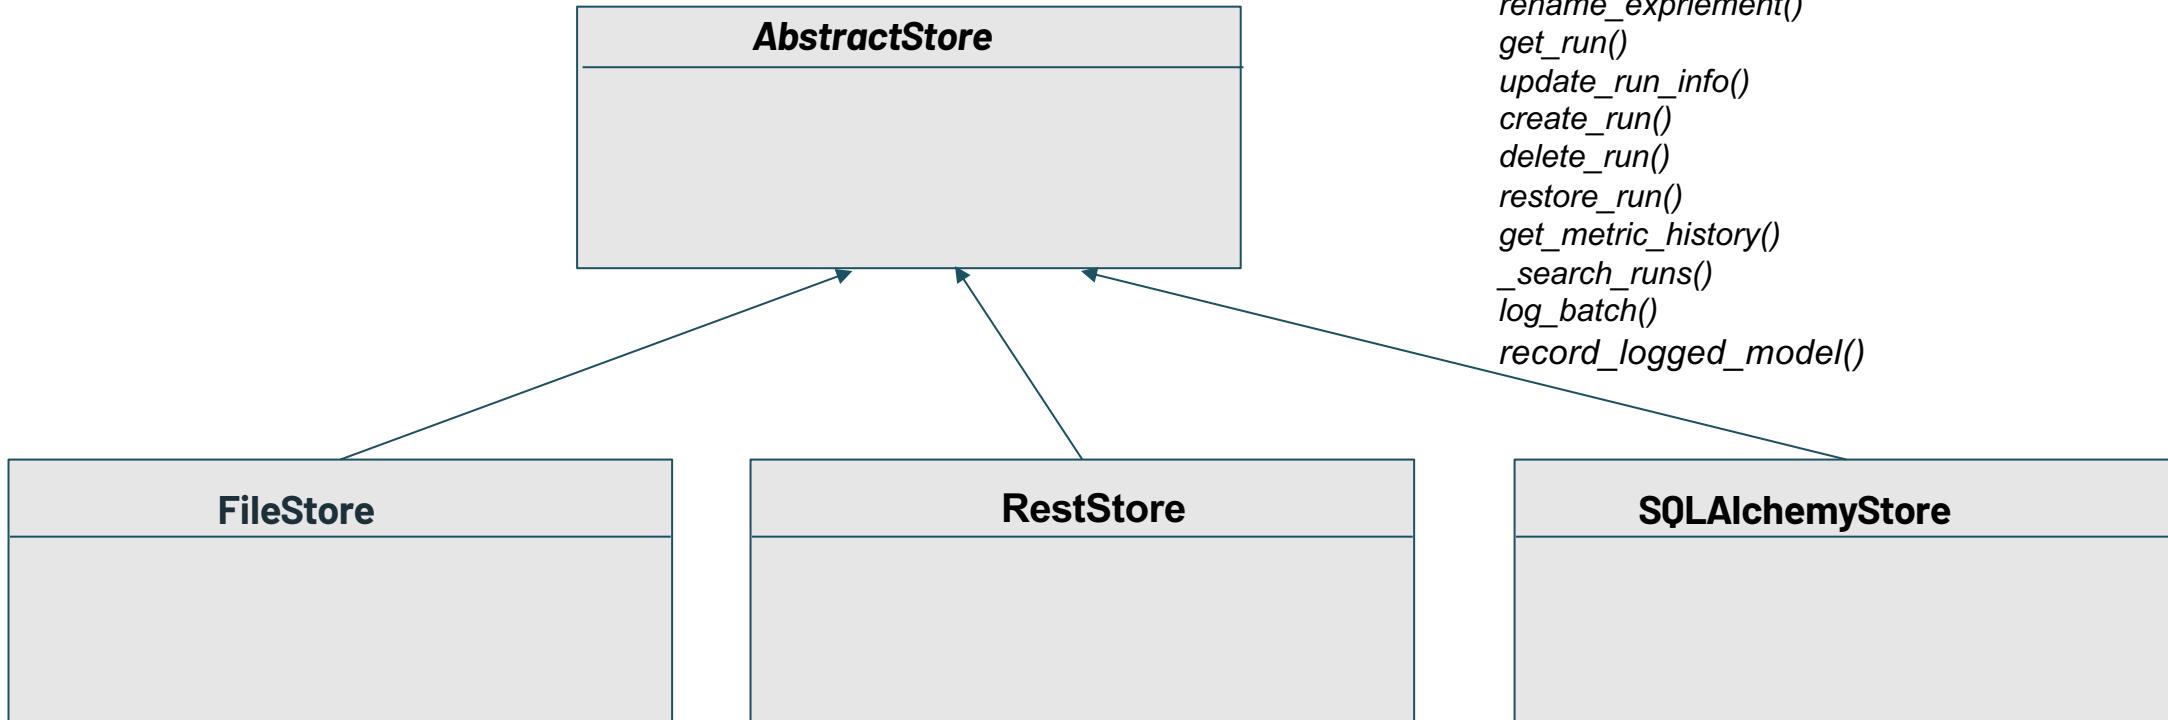
Entity (Metadata) Backend Store

- FileStore (local filesystem)
 - ***mlruns*** directory by default
- SQLStore (via SQLAlchemy)
 - PostgreSQL, MySQL, SQLite
- MLflow Plugins Scheme
 - Customized Entity Metastore
- Managed MLflow on Databricks
 - MySQL on AWS and Azure

Artifact Store

- Local Filesystem
 - ***mlruns*** directory
- S3 backed store
- Azure Blob storage
- Google Cloud Storage
- DBFS artifact repo

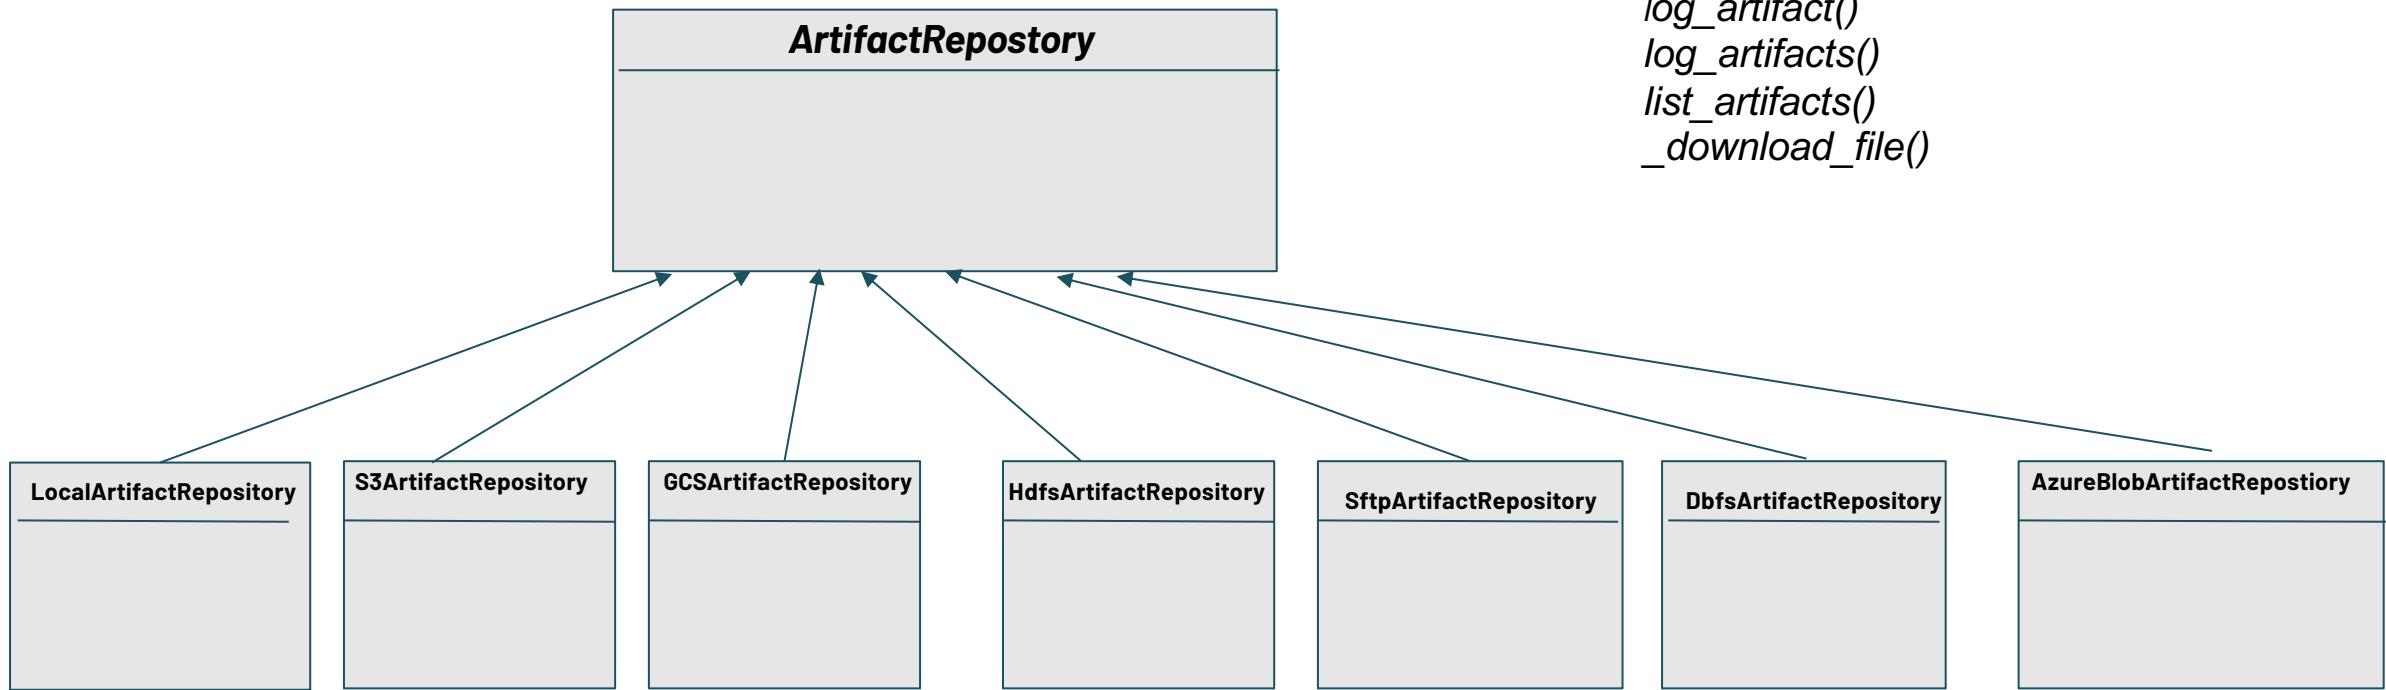
Backend Store Class Hierarchy



Abstract methods:

*list_experiments()
create_experiment()
delete_experiment()
restore_experiment()
rename_experiment()
get_run()
update_run_info()
create_run()
delete_run()
restore_run()
get_metric_history()
search_runs()
log_batch()
record_logged_model()*

ArtifactRepository Class Hierarchy



Abstract methods:

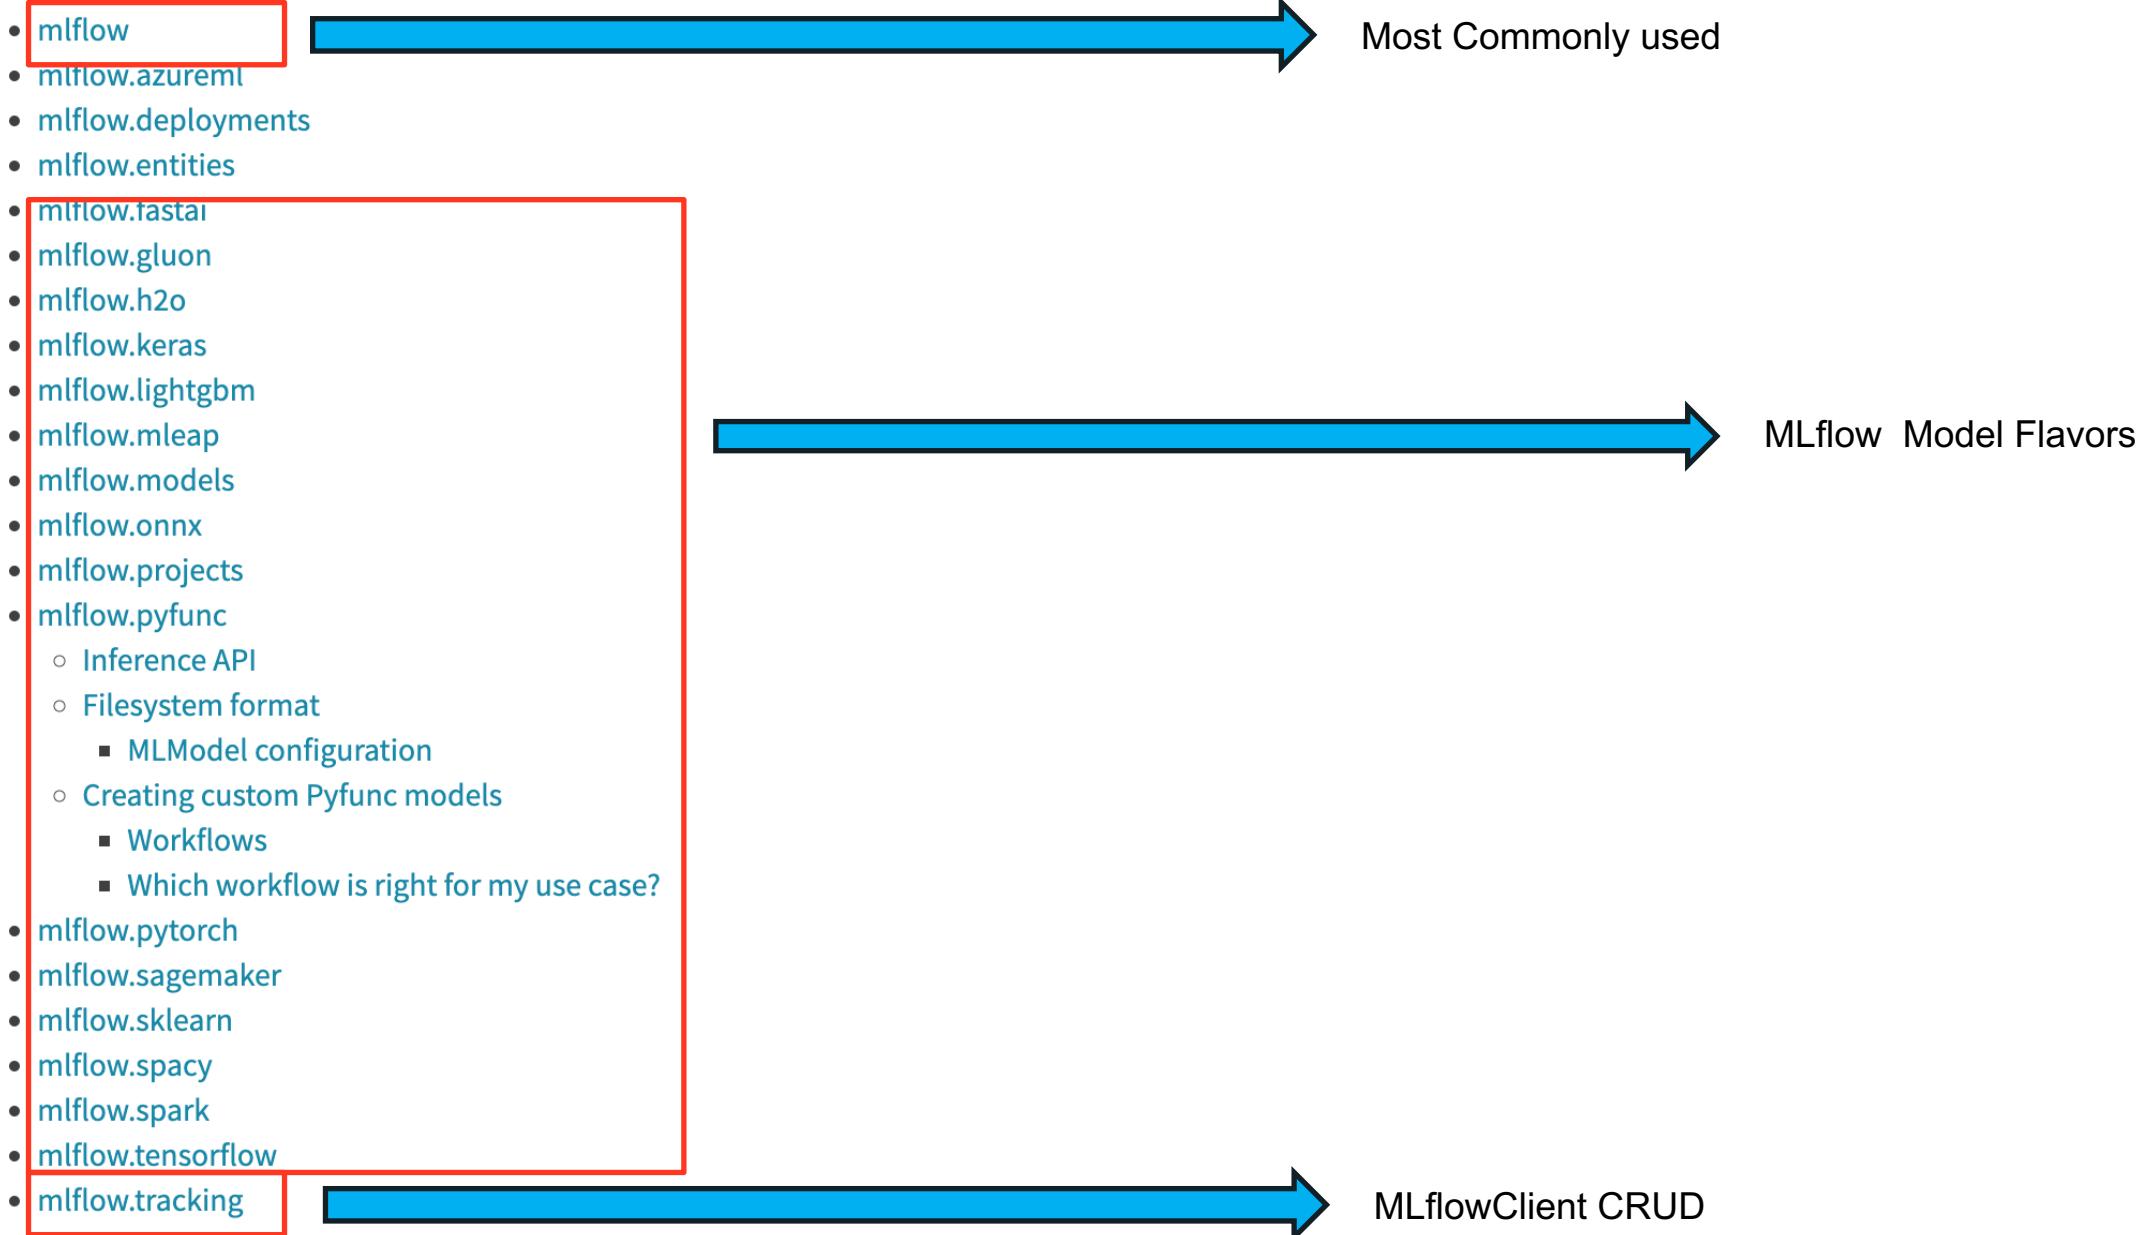
*log_artifact()
log_artifacts()
list_artifacts()
_download_file()*

Other Concrete classes include:

- DatabricksArtifactRepository
- DatabricksModelsArtifactRepository
- FtpArtifactRepository
- ModelsArtifactRepository
- RunsArtifactRepository

Python API

The MLflow Python API is organized into the following modules. The most common functions are exposed in the `mlflow` module, so we recommend starting there.



Built-In Model Flavors

MLflow provides several standard flavors that might be useful in your applications. Specifically, many of its deployment tools support these flavors, so you can export your own model in one of these flavors to benefit from all these tools:

- [Python Function \(python_function\)](#)
- [R Function \(crate\)](#)
- [H2O \(h2o\)](#)
- [Keras \(keras\)](#)
- [MLeap \(mleap\)](#)
- [PyTorch \(pytorch\)](#)
- [Scikit-learn \(sklearn\)](#)
- [Spark MLlib \(spark\)](#)
- [TensorFlow \(tensorflow\)](#)
- [ONNX \(onnx\)](#)
- [MXNet Gluon \(gluon\)](#)
- [XGBoost \(xgboost\)](#)
- [LightGBM \(lightgbm\)](#)



mlflow.sklearn

The `mlflow.sklearn` module provides an API for logging and loading scikit-learn models. This module exports scikit-learn models with the following flavors:

Python (native) pickle format

This is the main flavor that can be loaded back into scikit-learn.

mlflow.pyfunc

Produced for use by generic pyfunc-based deployment tools and batch inference.

mlflow.sklearn.get_default_conda_env(include_cloudpickle=False) [source]

Returns

The default Conda environment for MLflow Models produced by calls to `save_model()` and `log_model()`.

mlflow.sklearn.load_model(model_uri) [source]

Load a scikit-learn model from a local file or a run.

Parameters

model_uri -

The location, in URI format, of the MLflow model, for example:

- `/Users/me/path/to/local/model`
- `relative/path/to/local/model`

MLflow API Data Flow: Localhost defaults

- **No** mlflow server --backend-store-uri --default-root-artifact-uri
- **Everything** stored under local directory **mlruns**

```
with mlflow.start_run(run_name="DEFAULTS") as run:  
    params = {"n_estimators": 3, "random_state": 42}  
    sk_learn_rfr = RandomForestRegressor(params)  
  
    # Log params using the MLflow sklearn APIs  
    mlflow.log_params(params)  
    mlflow.log_param("param_1", randint(0, 100))  
    mlflow.log_metric("metric_1", random())  
    mlflow.log_metric("metric_2", random() + 1)  
    mlflow.sklearn.log_model(sk_learn_rfr,  
                            artifact_path="sklearn-model")
```



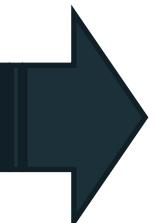
```
mlruns  
└── 0  
    ├── e7d505e713f14d64a74279f6c1a5d8e7  
    │   ├── artifacts  
    │   │   └── sklearn-model  
    │   │       ├── MLmodel  
    │   │       └── conda.yaml  
    │   └── model.pkl  
    ├── meta.yaml  
    ├── metrics  
    │   ├── metric_1  
    │   └── metric_2  
    ├── params  
    │   ├── n_estimators  
    │   ├── param_1  
    │   └── random_state  
    └── tags  
        ├── mlflow.log-model.history  
        ├── mlflow.runName  
        ├── mlflow.source.git.commit  
        ├── mlflow.source.name  
        ├── mlflow.source.type  
        └── mlflow.user  
    └── meta.yaml  
  
7 directories, 16 files
```

MLflow API Data Flow: Localhost w/ SQLAlchemy backend store

- No `mlflow server --backend-store-uri --default-root-artifact-uri`
- Artifacts stored under local directory `mlruns`
- MLflow entities stored in `mlruns.db`

```
local_store_uri = "sqlite:///mlruns.db"  
mlflow.set_tracking_uri(local_store_uri)
```

```
with mlflow.start_run(run_name="LOCAL_REGISTRY") as run:  
    params = {"n_estimators": 3, "random_state": 42}  
    sk_learn_rfr = RandomForestRegressor(params)  
  
    # Log params using the MLflow sklearn APIs  
    mlflow.log_params(params)  
    mlflow.log_param("param_1", randint(0, 100))  
    mlflow.log_metric("metric_1", random())  
    mlflow.log_metric("metric_2", random() + 1)  
    mlflow.sklearn.log_model(sk_learn_rfr,  
        artifact_path="sklearn-model")
```



```
mlruns  
└── 0  
    └── 1f896ddc61d34d089d4126eddceb5b52  
        └── artifacts  
            └── sklearn-model  
                ├── MLmodel  
                ├── conda.yaml  
                └── model.pkl
```

4 directories, 3 files

```
(tutorials) ➔ mlflow_fluent git:(master) ✘ sqlite3 mlruns.db  
SQLite version 3.30.1 2019-10-10 20:19:45  
Enter ".help" for usage hints.  
sqlite> .tables  
alembic_version      metrics          registered_model_tags  
experiment_tags      model_version_tags registered_models  
experiments          model_versions   runs  
latest_metrics       params           tags  
sqlite> select * from params;  
n_estimators|3|29c1e723cfad48bbbc20c4ca52414780  
random_state|42|29c1e723cfad48bbbc20c4ca52414780  
param_1|10|29c1e723cfad48bbbc20c4ca52414780  
sqlite> select * from metrics;  
metric_1|0.216184107170305|1597259612673|29c1e723cfad48bbbc20c4ca52414780|0|0  
metric_2|1.17106399018329|1597259612697|29c1e723cfad48bbbc20c4ca52414780|0|0  
sqlite> |
```

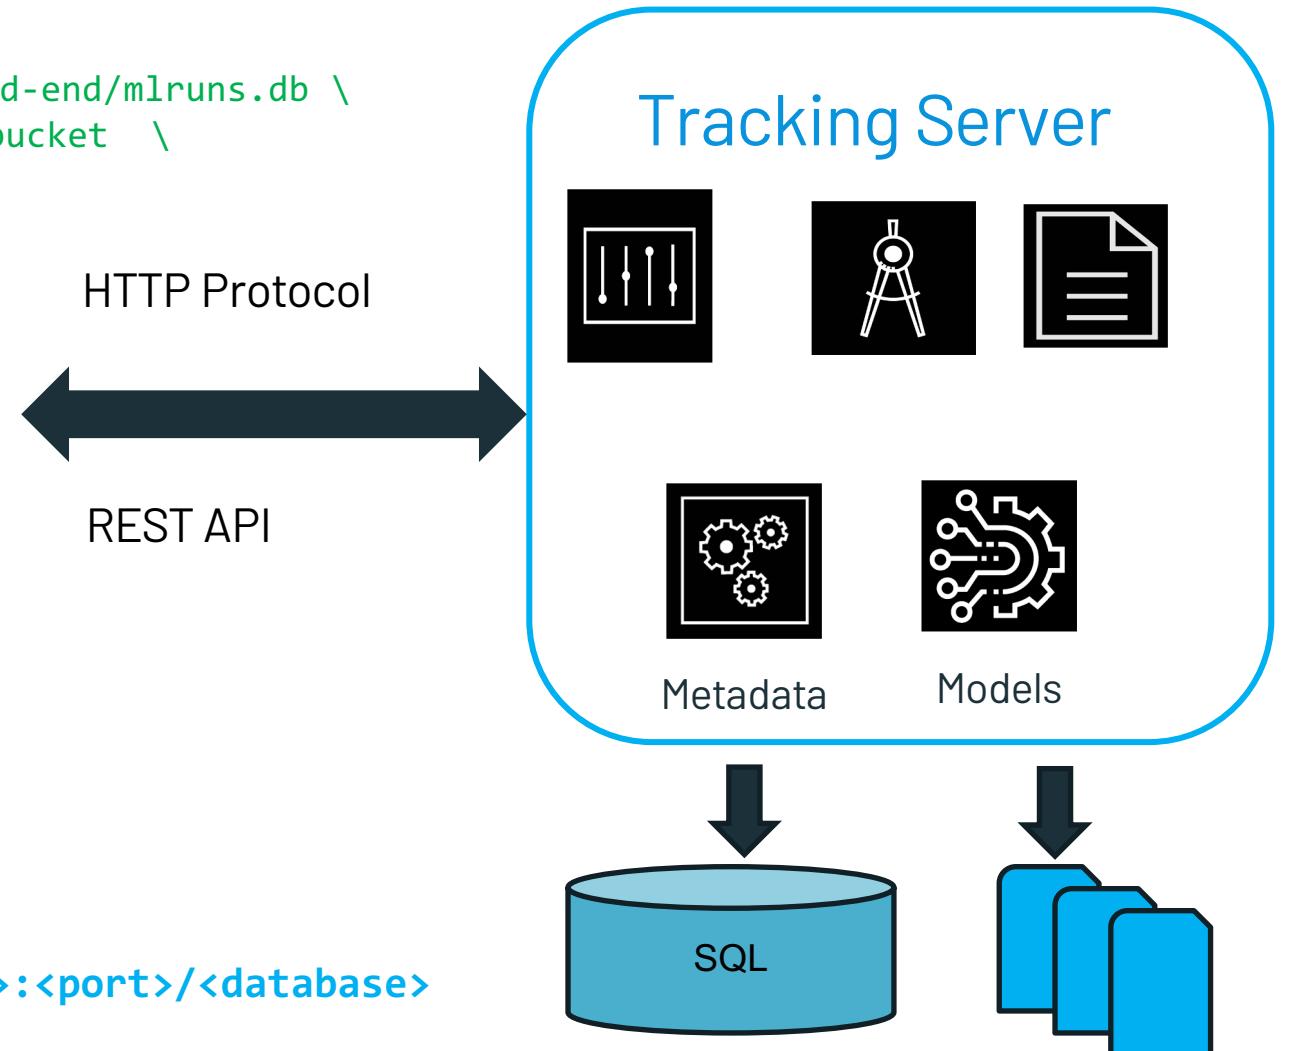
MLflow API Data Flow: Use Tracking Server SQLAlchemy backend store

```
mlflow server --backend-store-uri sqlite:///mnt/backend-end/mlruns.db \  
--default-artifact-root s3://my-mlflow-bucket \  
--host host.mydomain.com
```

```
tracking_uri = "https://host.mydomain.com:5000"  
mlflow.set_tracking_uri(tracking_uri)  
  
with mlflow.start_run() as run:  
    params = {"n_estimators": 3, "random_state": 42}  
    sk_learn_rfr = RandomForestRegressor(params)  
  
    # Log params using the MLflow sklearn APIs  
    mlflow.log_params(params)  
    mlflow.log_param("param_1", randint(0, 100))  
    mlflow.log_metric("metric_1", random())  
    mlflow.log_metric("metric_2", random() + 1)  
    mlflow.sklearn.log_model(sk_learn_rfr,  
        artifact_path="sklearn-model")
```

Backend Store URI Format: MySQL, PostgreSQL, SQLite, etc.

<dialect>+<driver>://<username>:<password>@<host>:<port>/<database>



MLflow set of APIs

Fluent MLflow APIs

- Python
 - High-level operations for runs and experiments
 - Model Flavor APIs
- Java
 - MLflowContext
 - Experiments, runs, search, etc
- R
 - Experiments, runs, search etc

MLflowClient

- Low Level CRUD interface to experiments, runs, Entities, Model Registry, etc
- `import mlflow.tracking`
- `client = MLflowClient(**kwargs)`
- Metric, Param, Search etc

MLflow REST API

- Make REST calls to Tracking server with endpoints
- `https://<tracking_server>/api/..`
- `/2.0/mlflow/2.0/runs/log-metric`
- `/2.0/mlflow/runs/log-parameter`
- `/2.0/mlflow/runs/search`

mlflow Tracking for Model Schemas

Record what fields are consumed & produced by the model to prevent data mismatches

in MLflow 1.11

```
with mlflow.start_run(run_name='keras'):
    # log model and datasource
    mlflow.keras.autolog()
    mlflow.spark.autolog()

    sig = infer_signature(X_train, y_train)
```

Schema	
Name	Type
Inputs (7)	
user_purchases_7d	long
user_purchases_14d	long
user_location	string
user_last_login	string
item_inventory	long
Outputs (1)	
prediction	double

mlflow Tracking for Interpretability

SHAP library feature importances and visualizations

in MLflow 1.12

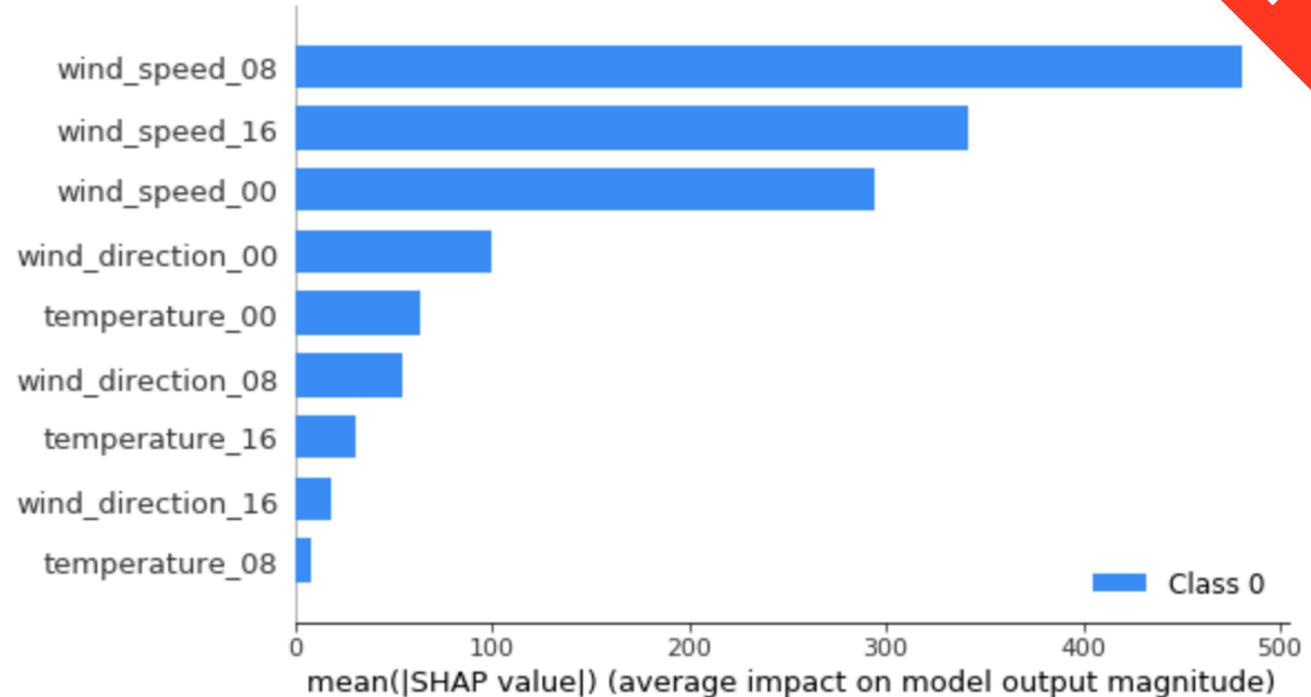
```
with mlflow.start_run(run_name='keras'):
    # log model and datasource
    mlflow.keras.autolog()
    mlflow.spark.autolog()

    sig = infer_signature(X_train, y_train)

    mlflow.shap.log_explanation(model, X_train[:100])
```



SHAP



Easily log and query all aspects of your ML development process

mlflow + PyTorch

in MLflow 1.12

Collaboration between Facebook and Databricks to bring ML platform features to PyTorch

MLflow autologging for PyTorch Lightning

TorchScript support for faster packaged models

Model deployment to TorchServe

MLflow 1.12 Features Extended PyTorch Integration

And model explainability, autologging, and other enhancements



by Jules Damji, Siddharth Murching and Harutaka Kawamura
Posted in ENGINEERING BLOG | November 13, 2020

mlflow + PyTorch

```
# Enable MLflow autologging
import mlflow.pytorch_lightning
mlflow.pytorch_lightning.autolog()

# Train a model with PyTorch Lightning
model = MyModel()
trainer = pl.Trainer()
trainer.fit(model, train, val)
```

```
# Deploy to TorchServe
mlflow deployments create -t torchserve -m models:/mymodel/production -n mymodel
```

mnist_autolog_log_niter > Run b288b477bf3f4ab18c52a8ef26e4c118 ▾

Date: 2020-07-27 11:02:09 Source: __main__.py User: ec2-user

Duration: 2.3min Status: FINISHED

▼ Notes [🔗](#)

None

▼ Parameters

Name	Value
epsilon	1e-08
learning_rate	0.001
optimizer_name	Adam

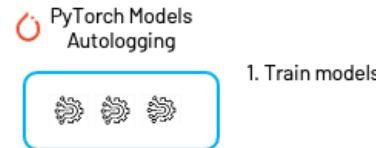
▼ Metrics

Name	Value
avg_test_acc 🔗	0.968
avg_train_loss 🔗	0.029
avg_val_loss 🔗	0.113
loss 🔗	0.003

▼ Tags

Name	Value	Actions

mlflow + PyTorch



MLflow and PyTorch — Where Cutting Edge AI meets MLOps



Authors: Geeta Chauhan, PyTorch Partner Engineering Lead and Joe Spisak, PyTorch Product Lead at Facebook

in MLflow 1.13

New Tracking APIs

Extend Tracking APIs for new types of Artifacts

- `mlflow.log_text("text1", "file1.txt")`
- `mlflow.log_text("<h1>header</h1>", "index.html")`
- `mlflow.log_dict(dictionary, "dict.json")`
- `mlflow.log_figure(figure, "figure.png")`
 - [matplotlib.figure.Figure](#)
 - [plotly.graph_objects.Figure](#)
- `mlflow.log_image(image, "image.png")`
 - [PIL.Image.Image](#)

```
1 import mlflow
2
3 with mlflow.start_run():
4     # Log text to a file under the run's root artifact
5     # directory
6     mlflow.log_text("text1", "file1.txt")
7
8     # Log text in a subdirectory of the run's root artifact
9     # directory
10    mlflow.log_text("text2", "dir/file2.txt")
11
12    # Log HTML text
13    mlflow.log_text("<h1>header</h1>", "index.html")
```

What Did We Talk About? **mlflow**TM

- Modular Components greatly simplify the ML lifecycle
- Easy to install and use & Great Developer Experience
- Develop & Deploy locally and track locally or remotely
- Available APIs: Python, Java & R (Soon Scala)
- Visualize experiments and compare runs

Learning More About MLflow & Get Involved!

- `pip install mlflow` – to get started
- Find docs & examples at mlflow.org
- Peruse code and contribute at [MLflow Github](https://github.com/mlflow/mlflow)
- Join the [Slack channel](#)
- More [MLflow tutorials](#)

MLflow Tracking Tutorials

<https://github.com/dmatrix/olt-mlflow>

Thank you! 😊

Q & A

jules@databricks.com
@2twitme

<https://www.linkedin.com/in/dmatrix/>